



AWS Lambda

Beginner to Expert

HELLO!

I'm Daniel

Senior Software Engineer

youtube.com/c/beabetterdev





COURSE INTRODUCTION

Approach, Teaching Format, Content, Pre-requisites



Course Introduction

- ◊ Teach you the fundamentals of AWS Lambda
 - Emphasis on the important bits
- ◊ Visual Learning
 - Console Walkthroughs
 - Cloud Development Kit for Infrastructure as Code



Course Contents

Introduction

- What is a Lambda Function?
- A Brief Tour of the Console
- How does it work?

Good to Knows

- IaC with CDK
- Pricing

Key Concepts

- Configuration
- Deployments
- Monitoring

Service Integrations

- Event Processing & Workflows
- APIs & Adhoc Jobs
- Timed Jobs

Advanced Concepts

- Invocation Types
- Throttling & Concurrency
- Performance Tuning



Who Are You?

- ◊ Developers
 - ◊ Data Engineers & Scientists
 - ◊ Devops
 - ◊ Other Industry Professional
- Folks starting a new job and looking to ramp up quickly**



Pre-requisites

- ❖ **Python Fundamentals**
 - Functions, variables primitives, etc...
- ❖ **General Understanding of AWS**
- ❖ **Python, PIP, AWS CLI, CDK**
- ❖ **AWS Account**



The Value of Learning Lambda

- ❖ Massive Industry Adoption
- ❖ Applications in Many Job Families
- ❖ Extensive Community Support

Lower Cost vs Traditional Bare Metal



Introduction to AWS Lambda

Agenda

- ❖ What is AWS Lambda?
- ❖ Why is Lambda Useful?
- ❖ Disadvantages
- ❖ Who Uses Lambda?



History of AWS Lambda

Early 2000s

Traditional Company Specific
Data Centers



Mid to Late 2000s

Early Cloud Infrastructure
(EC2)



2014+

Enhanced
Cloud Infrastructure



What is AWS Lambda?

- ◆ Compute Service
- ◆ Run code at scale without worrying about servers
- ◆ You write *Functions* - the primary unit of Lambda
- ◆ Useful in many applications such as
 - ◆ API Hosting
 - ◆ Event Processing
 - ◆ Ad-hoc or timer based jobs



A Typical Workflow

1

Create a Function

2

Write & Upload Your Code

3

Run Your Function



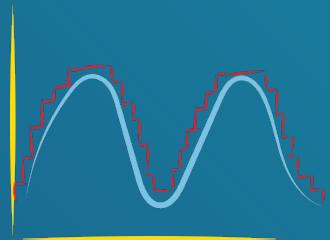


Why is Lambda Useful?

No Servers to Manage



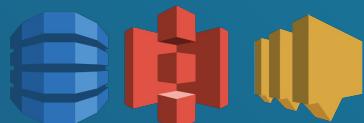
Autoscaling



Performance



Service Integrations



Pay For What You Use



Easy To Use

“

“With AWS Lambda, you gain flexibility at the expense of control.”



Users of AWS Lambda

The Coca-Cola Company



NETFLIX



Udemy

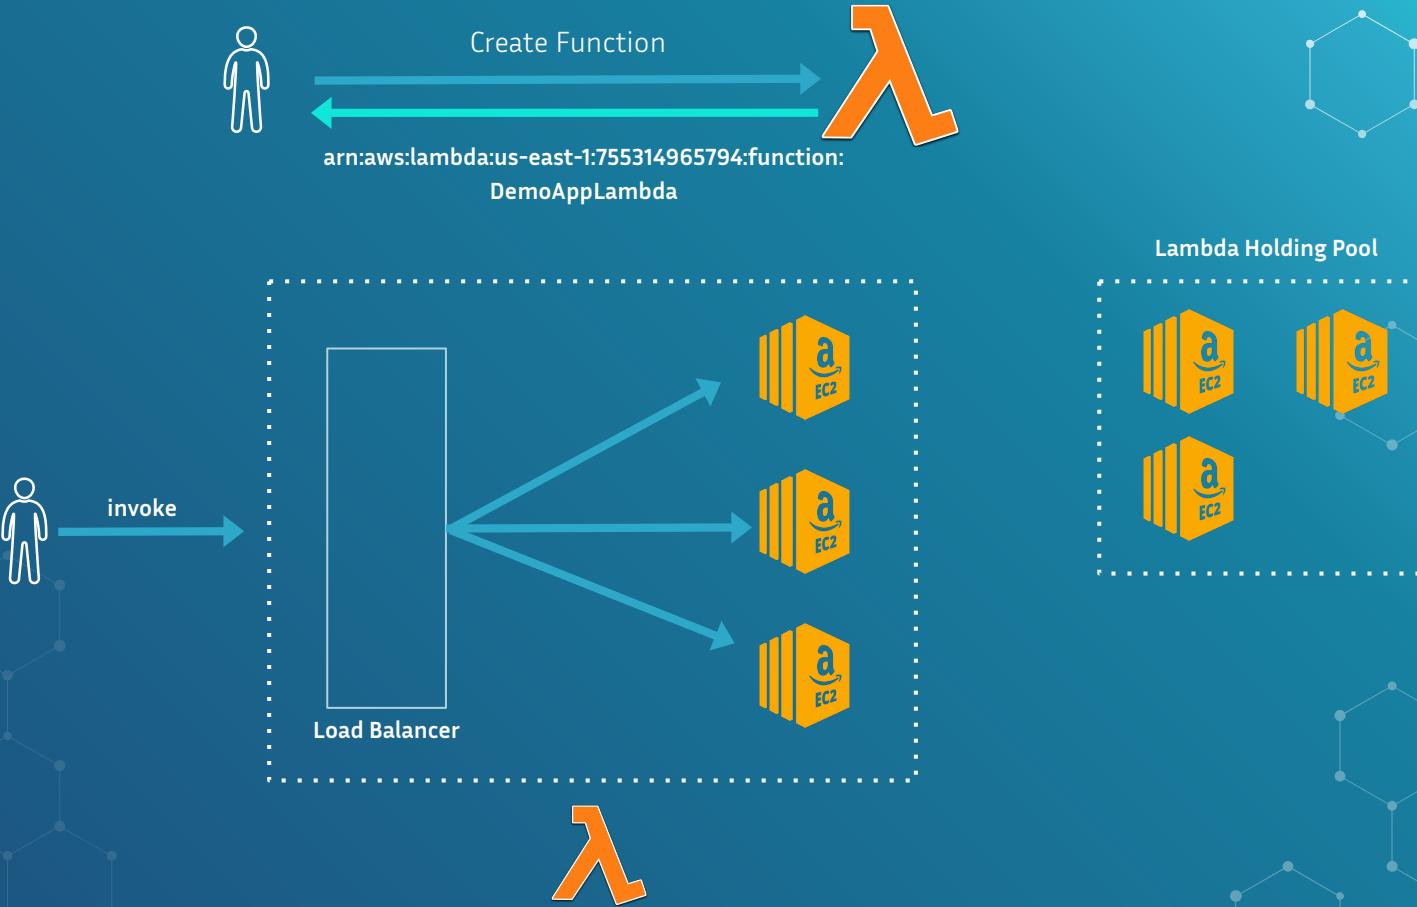


Robot



How Lambda Works Under The Hood

Under The Hood





Core Concepts - Triggers

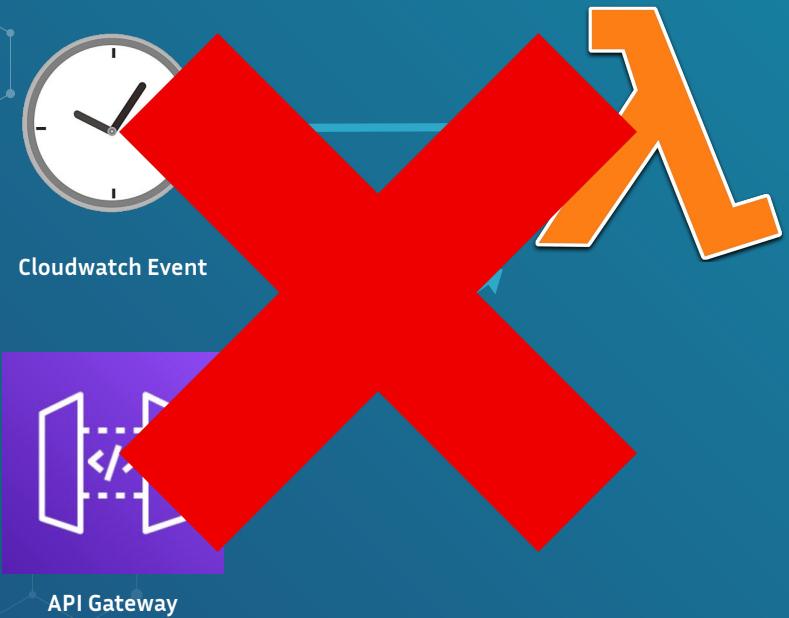
Not like this...



What is a Trigger?

- ❖ Means to Invoke your Function
- ❖ SDK, CLI, Console, Service Integration
- ❖ Can have many triggers
- ❖ Similar concept called *Event Source Mapping*
- ❖ Different Triggers have different scaling / retry behaviour

Don't Try to Do This



Common Triggers / Event Source Mappings

- ◆ **Cloudwatch Events** for Timer based jobs
- ◆ **API Gateway** for REST APIs
- ◆ **SNS** or **SQS** for event processing
- ◆ **DynamoDB Streams** for Record Changes
- ◆ **S3** for Object Creation / Updates



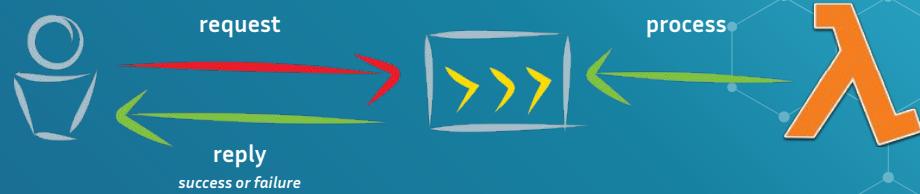
Core Concepts - Invocation Types

Synchronous and Asynchronous

Synchronous



Asynchronous



How do we handle failure?

Dead Letter Queue

Why does this matter?

- ◊ Synchronous can have cost implications
- ◊ Asynchronous can be more suitable for fault tolerant use cases
- ◊ Different Failure Behaviours

Real Life Example

- ◊ Event Processing System With Multiple Lambdas calling each other
- ◊ Initially Used Synchronous... Whoops
- ◊ Switched to Async and Voila...





Core Concepts - Duration / Timeouts



We all have our Limits

Lambda Duration Limits

- ◆ Maximum 15 minutes before automatic timeout
- ◆ This means Lambda are not suitable for long running processes
- ◆ Lambdas will automatically Timeout when the configured limit is reached
- ◆ Default is 3 seconds - Configurable via Console

Pro Tips

- ◆ Longer timeouts aren't always better
- ◆ Background (asynchronous) processes are also terminated



Core Concepts - Memory

Who Remembers These?



Lambda Memory

- ◆ Defines the amount of memory available to your function at runtime
- ◆ Range between 128 MB and 10,240 MB
- ◆ Changing Memory is the most significant performance lever available to you
- ◆ Increasing memory determines instance type behind the scenes

Pro Tips

- ❖ Increasing memory setting can reduce the impact of cold start
- ❖ Can also help with Networking Connectivity Speed
- ❖ Memory increase doesn't result in linear performance gain
- ❖ You should try and find the sweet spot to balance performance and invocation duration



Side Topic - Pricing

Invocation Count

The number of times you invoke the function
\$.20 per 1 Million Requests

Memory

The amount of memory you provision

Memory (MB)	Price per 1ms
128	\$0.000000021
512	\$0.000000083
1024	\$0.000000167
1536	\$0.000000250
2048	\$0.000000333
3072	\$0.000000500
4096	\$0.000000667
5120	\$0.000000833
6144	\$0.000001000
7168	\$0.000001167
8192	\$0.000001333
9216	\$0.000001500
10240	\$0.000001667

Duration

How long your invocations run for
\$.0000166667 per GB Second

3 Million Invocations

512 MB Memory

1 Second Duration

\$18.74 / month

30 Million Invocations

128 MB Memory

200ms Duration

\$11.63 / month

To Estimate Your Costs

- ❖ Use the Cost Calculator!
- ❖ <https://aws.amazon.com/lambda/pricing/>

Pro Tips

- ◊ Start with low memory settings
- ◊ Once you collect enough invocation history, use the Compute Optimizer Tool
- ◊ 1 Million free requests per month and 400,000 GB Seconds



Core Concepts - Function Execution & Cold Start

Decomposing a Function Execution

Code Download

1

Start Execution Environment

2

Full Cold Start

Execute Init Code

3

Partial Cold Start

Execute Handler Code

4

Warm Start

```
1 //1 - Import required libraries
2 var AWSXRay = require('aws-xray-sdk-core')
3 var captureMySQL = require('aws-xray-sdk-mysql')
4 var mysql = captureMySQL(require('mysql2'))
5 var Sequelize = require('sequelize')
6
7 //2 - Extract environment variables from runtime
8 const username = process.env.databaseUser
9 const password = process.env.databasePassword
10 const host = process.env.databaseHost
11
12 //Function Entry Point
13 exports.handler = async (event) => {
14     //3 - Create Database Connection
```



Cold Start also occurs while scaling up

Strategies to Minimize Cold Start

- ❖ Minimize number of library dependencies
 - ❖ Only import what you need
 - ❖ Raise Memory Configuration
- ❖ Utilize Provisioned Concurrency

Example - Querying a Database

```
1 //1 - Import required libraries
2 var AWSXRay = require('aws-xray-sdk-core')
3 var captureMySQL = require('aws-xray-sdk-mysql')
4 var mysql = captureMySQL(require('mysql2'))
5 var Sequelize = require('sequelize')
6
7 //2 - Extract environment variables from runtime
8 const username = process.env.databaseUser
9 const password = process.env.databasePassword
10 const host = process.env.databaseHost
11
12 //Function Entry Point
13 exports.handler = async (event) => {
14     //3 - Create Database Connection
15     var connection = mysql.createConnection({ //
16         host      : host,
17         user      : username,
18         password : password,
19         database : 'lambdadb'
20     });
21
22     //4 - Define Query
23     var query = "SELECT * FROM Customers";
24     var result;
25
26     //5 - Perform Query Using Connection
27     connection.query(query, function (error, results, fields) {
28         if (error) throw error
29         console.log("Ran query: " + query)
30         for (result in results)
31             console.log(results[result])
32     })
33 }
```



Core Concepts - Concurrency & Throttling



32
BẮC GIÁP BẮT → NHƠN

XÃ BUS CỦ MƯỜU



What is Lambda Concurrency?

- ❖ # of requests being served at a given moment
- ❖ New containers are spawned for each concurrent request
- ❖ Concurrency is a major scaling consideration, and can cause applications to fail due to *Throttling*
- ❖ Default 1000 units of concurrency per AWS account per region
- ❖ Unreserved, Reserved, and Provisioned

Throttling aka *RateExceeded*

- ◆ Throttling is when Lambda rejects a request
- ◆ Occurs when in flight invocations exceeds available concurrency

Unreserved

Common concurrency pool

Reserved

Dedicated concurrency pool

Provisioned

Dedicated and 'always on' concurrency pool



Core Concepts - Reserved & Unreserved Concurrency

Unreserved Concurrency

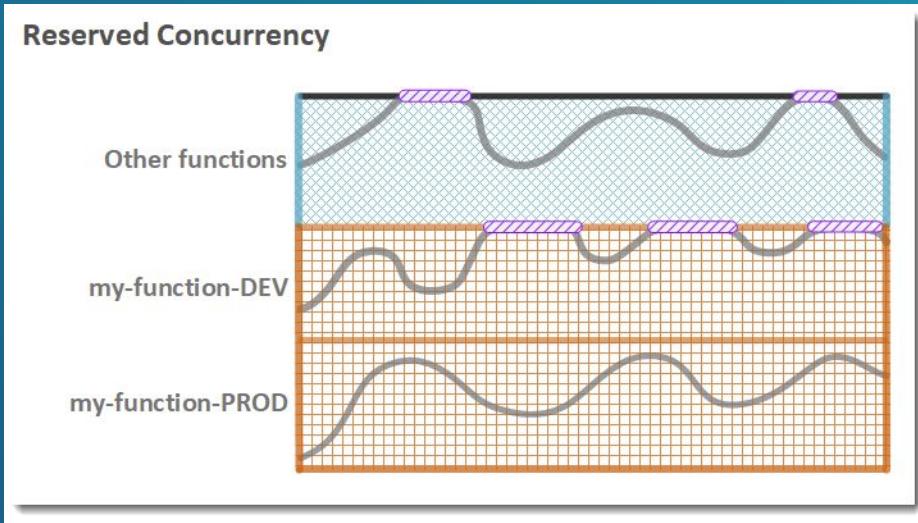
- ◆ Free for all among all functions in an account per region
- ◆ If one function consumes all concurrency, the others will get throttled
- ◆ Limit can be raised via AWS Support Ticket



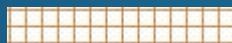
Reserved Concurrency

- ◆ A portion is always reserved for a specific function, even if there are no invocations
- ◆ You can use reserved concurrency to minimize or maximize your processing rate
- ◆ Good for predictable resource allocation
- ◆ Bad in terms of utilization of resources

Reserved



Function Concurrency



Reserved Concurrency



Unreserved Concurrency



Throttling

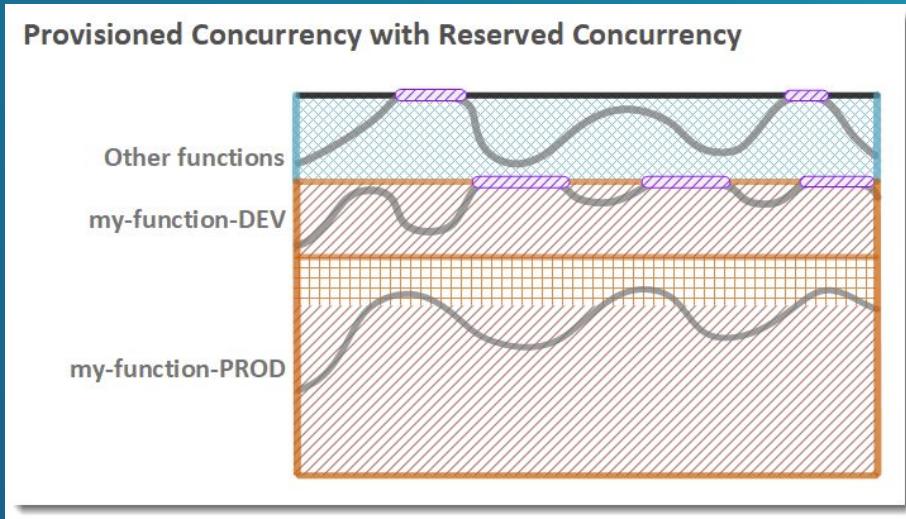


Core Concepts - Provisioned Concurrency

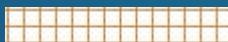
Provisioned Concurrency

- ◆ Pool of concurrency that is 'always on'
- ◆ Solves the Cold Start problem, *kinda*
- ◆ Supports autoscaling policies based on usage
- ◆ Very expensive (34\$/month for 5 units always on)
- ◆ Blurs the line between serverless and metal

Provisioned with Reserved



Function Concurrency



Reserved Concurrency



Provisioned Concurrency



Unreserved Concurrency



Throttling

Pro Tips

- ◊ Alarm on throttles for early indicators of issues
 - ◊ Evaluate your concurrency needs and plan accordingly
 - ◊ Have your clients use Exponential Backoff to avoid retry storms
- Raising your Memory Limit can help, but *be careful*
- ◊ Use a small amount of provisioned concurrency to mitigate cold starts for latency sensitive apps



Core Concepts - Versions

charge

Versions

- ◆ Versions are snapshots of the code and settings of a function.
 - ◆ Code and settings get locked to a version on publish
 - ◆ Allow you to have multiple 'versions' of a function at the same time - useful for quick rollbacks
- ◆ New uploads default to \$LATEST version

`arn:aws:lambda:us-east-1:755314965794:function:DemoAppLambda:1`

Versions - Example



V1



V2



V3



Version

1

\$LATEST



2

\$LATEST



3

\$LATEST



Core Concepts - Aliases

Aliases

- ◆ An Alias is a named pointer to a specific Version
`arn:aws:lambda:us-east-1:755314965794:function:DemoAppLambda:1`
- ◆ Useful for 'beta' testing, or 'pre-prod' environments

Aliases - Example



V1



V2



V3



Version

1

\$LATEST

prod



2

\$LATEST

prod



3

\$LATEST

prod

Advanced Concept - Version/Alias Weights

- ◊ You can assign weights to different versions using a Version/Alias
- ◊ *90% of Traffic to Version 1, 10% to Version 2*
- ◊ Useful for version validation prior to full blown deployment
- ◊ You need to use Aliases/Versions in order to use *Provisioned Concurrency*



Core Concepts - Environment Variables

Environment Variables

- ◊ Pairs of Strings with a key and value
- ◊ Allow you to adjust function behavior without touching your code
- ◊ *Examples*
 - ◊ STAGE
 - ◊ CLIENT_ENDPOINT_URL
 - ◊ DATABASE_CONNECTION_STRING

```
import os
db_connection_string = os.environ['DATABASE_CONNECTION_STRING']
region = os.environ['AWS_REGION']
```

An Example

TransactionDB-DEV



```
import os  
transaction_db_name = os.environ['TRANSACTION_DB_NAME']
```

TransactionProcessorLambda-DEV

DEVELOPMENT

TransactionDB-PROD



```
import os  
transaction_db_name = os.environ['TRANSACTION_DB_NAME']
```

TransactionProcessorLambda-PROD

PRODUCTION



Core Concepts - VPC & Networking

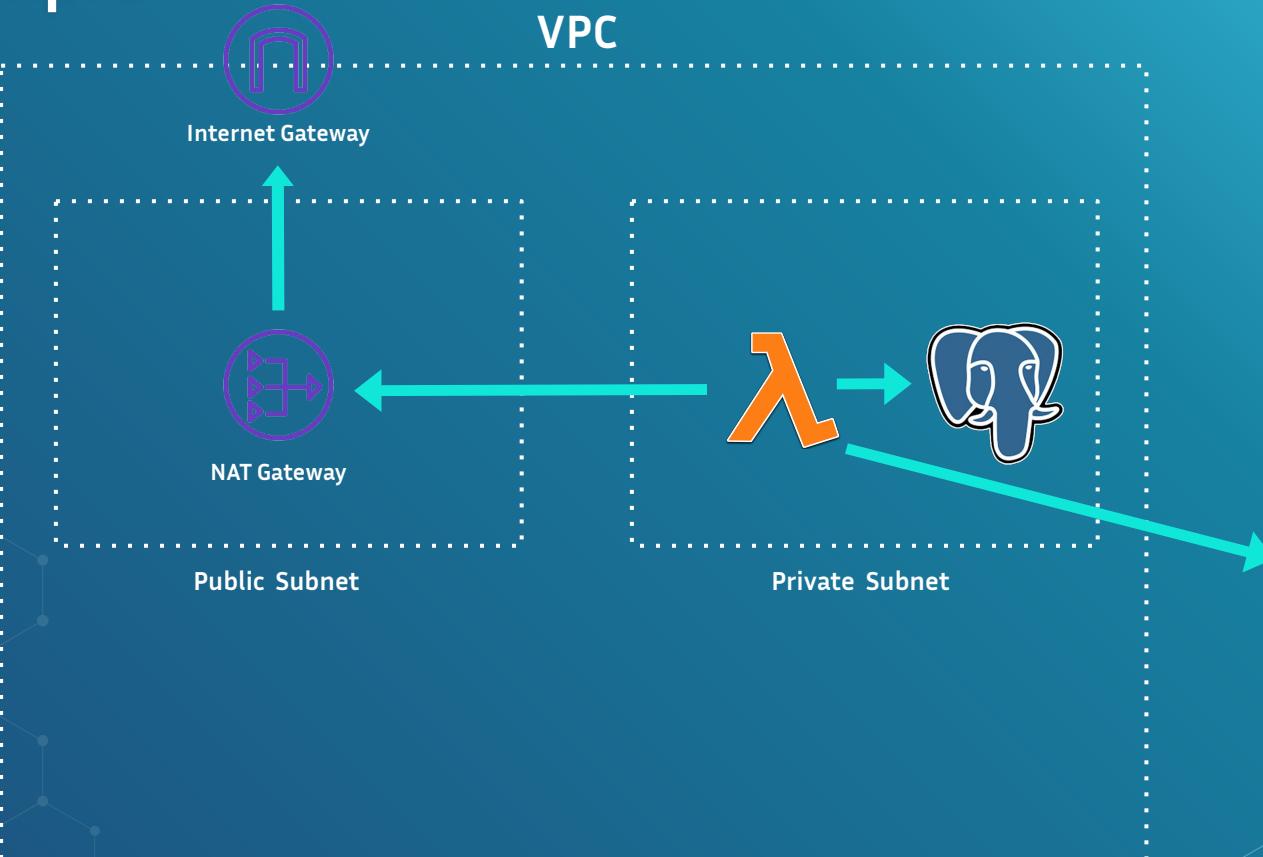
VPC - Quick Refresher

- ◆ Stands for Virtual Private Cloud
- ◆ Isolated and private network within AWS
- ◆ Public and Private Subnets
- ◆ *Critical for security and apps with compliance guidelines*

VPC and Lambda

- ◆ Only Necessary when your function needs to access resources not accessible over public internet
- ◆ Example: RDS or ElasticSearch in private subnet
- ◆ Behind the scenes, Lambda creates ENIs for each subnet the function is deployed into
- ◆ No longer contributes to cold start latency
- ◆ VPC Endpoints can be used to communicate with some AWS services privately

Example





You Specify:

VPC

Subnets

Security Groups

Lambda Creates:

Elastic Network Interface

Exposing Your Function to Other VPCs

- ❖ Use VPC Endpoints to privately connect to AWS Lambda
- ❖ Does not expose your function to the public internet
- ❖ You can limit *who* from the corresponding VPC can call your function

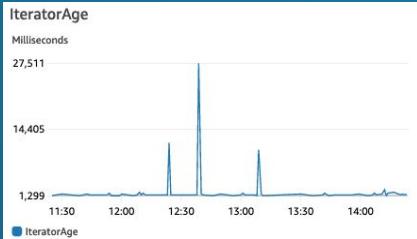
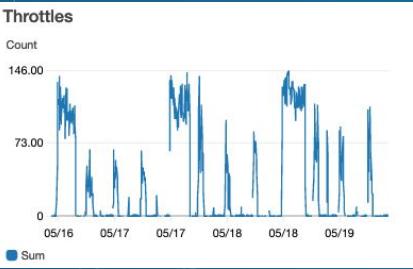
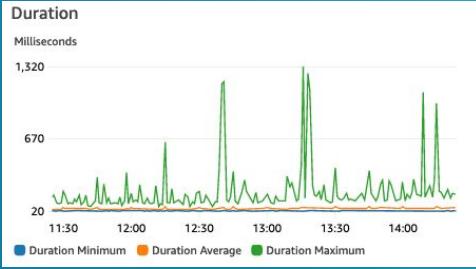
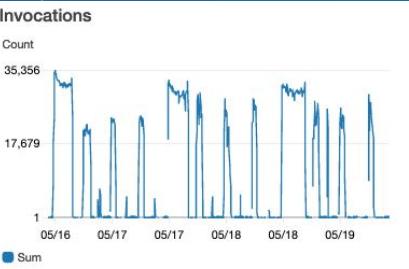
Good To Knows

- Launching in VPC can lead to longer startup times due to ENI connection overhead



Monitoring - Key Metrics

Metrics to Monitor



+ any other business metrics!

Important Tip

Metrics don't help if you're not looking
at them

SETUP ALARM ON YOUR METRICS!

Other Notable Metrics

- ◊ ConcurrentExecutions
- ◊ UnreserveredConcurrentExecutions



Monitoring - Logging

Logging

- Writing to stout automatically writes logs to Cloudwatch
- Log lines by default include metadata about the execution

```
▶ 2021-03-06T15:26:51.080-05:00      START RequestId: 1e379bd7-8723-4b12-ab14-efd904a317c2 Version: $LATEST
▶ 2021-03-06T15:26:53.712-05:00      * Serving Flask app "app" (lazy loading)
▶ 2021-03-06T15:26:53.712-05:00      * Environment: production
▶ 2021-03-06T15:26:53.712-05:00      WARNING: Do not use the development server in a production environment.
▶ 2021-03-06T15:26:53.712-05:00      Use a production WSGI server instead.
▶ 2021-03-06T15:26:53.712-05:00      * Debug mode: off
▶ 2021-03-06T15:26:53.772-05:00      * Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
▶ 2021-03-06T15:26:54.086-05:00      END RequestId: 1e379bd7-8723-4b12-ab14-efd904a317c2
▶ 2021-03-06T15:26:54.086-05:00      REPORT RequestId: 1e379bd7-8723-4b12-ab14-efd904a317c2 Duration: 3003.39 ms Billed Duration: 3000 ms Memory Size: 128 MB Max Memory Used: ...
▶ 2021-03-06T15:26:54.086-05:00      2021-03-06T20:26:54.086Z 1e379bd7-8723-4b12-ab14-efd904a317c2 Task timed out after 3.00 seconds
```



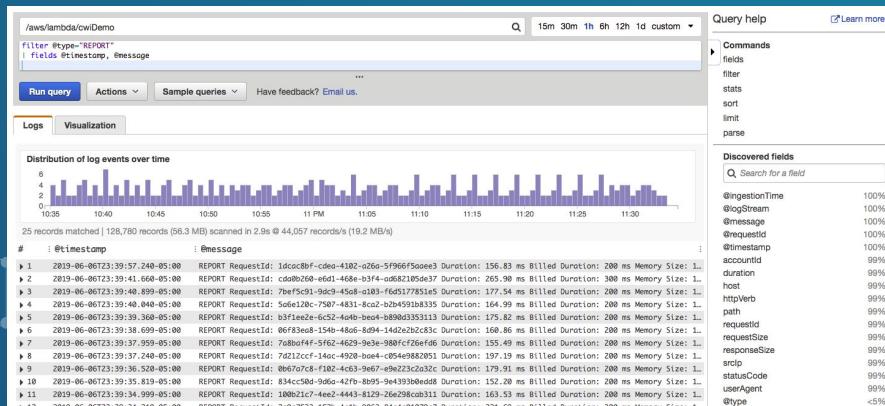
Metric Filters

- Allow you to write patterns (think regex) to extract metrics from log lines
- Much cheaper than using PutMetric
- More suitable for high TPS applications

*LATENCY - CreditCardServiceLatency 990ms
ERROR - InvalidArgumentException 1*

Cloudwatch Logs Insights

- ◆ Search for patterns across multiple functions at once
- ◆ Pay per use (amount of data scanned)
- ◆ Great 'Example Queries' that you can build from

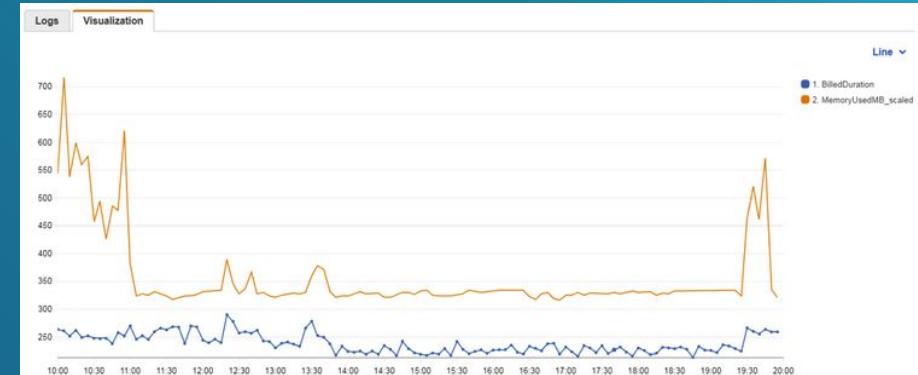


Cloudwatch Logs Insights Examples

```
fields (@maxMemoryUsed/1024/1024) as MemoryMB  
| filter @type="REPORT" and (MemoryMB>=62 and MemoryMB<=70)  
| stats count(@billedDuration) by MemoryMB, @billedDuration
```

#	MemoryMB	@billedDuration	count(@billedDuration)
1	69.6182	300	7
2	69.6182	400	21
3	69.6182	500	4
4	69.6182	200	16
5	68.6646	1000	1

```
filter @type="REPORT"  
| stats avg(@billedDuration) as BilledDuration,  
avg(@maxMemoryUsed/1024/1024) as MemoryUsedMB_scaled by bin(5m)
```





Pro Tips

- ◊ Don't be overly verbose with Logging
- ◊ Embed log lines with important ids to make tracing easier
- ◊ Use Cloudwatch Logs Insights for fuzzy searching
- ◊ Set up a log retention policy or archive regularly

```
[INFO] 2021-08-01 18:00:04 Lambda Execution Start
[INFO] 2021-08-01 18:00:05 PURCHASE Transaction with TransactionId=90a0fce1 and VALUE=10.00
[INFO] 2021-08-01 18:00:05 [TransactionId=90a0fce1] Processing PURCHASE
[INFO] 2021-08-01 18:00:06 [TransactionId=90a0fce1] Saved to DB
[INFO] 2021-08-01 18:00:07 [TransactionId=90a0fce1] Notified Subscribers via SNS
[INFO] 2021-08-01 18:00:08 [TransactionId=90a0fce1] Processing Complete
```



Monitoring - Lambda Insights



What is Lambda Insights?

- ◆ Performance Monitoring tool for Lambda
- ◆ Collects, aggregates & summarizes system level metrics (CPU, Memory, Network Usage)
- ◆ Embedded dashboard



How it works

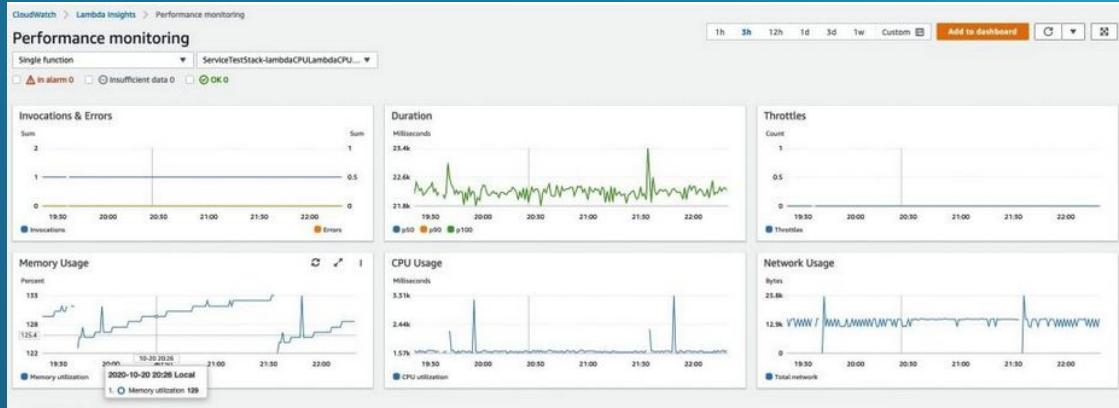
- ◆ Activating it adds a Lambda Extension to your function. Logs get emitted on every invocation.
- ◆ Metrics get sent to cloudwatch and extracted with metric filter patterns

Log Entry Example

2021-01-15T10:27:57.336-08:00

{"runtime": "Custom", "total_network": 84, "used_memory_avg":

```
{  
    "runtime": "Custom",  
    "total_network": 84,  
    "used_memory_avg": 112,  
    "request_id": "49d58d40-a0af-4cda-ab11-cf37ca27325d",  
    "agent_memory_avg": 7134549,  
    "total_memory": 128,  
    "cpu_total_time": 370,  
    "agent_memory_max": 7778304,  
    "used_memory_max": 125,  
    "fd_use": 384,  
    "tmp_max": 551346176,  
    "memory_utilization": 98,  
    "function_name": "Services-StepFnlambdaSteppriceLessThan556D8B304A-0LERCJ9Y1R7F",  
    "version": "$LATEST",  
    "agent_version": "1.0.54.0",  
    "trace_id": "1-6001de56-4f9b0bff0c00348c2f757547",  
    "threads_max": 10,  
    "event_type": "performance",  
    "fd_max": 18781,  
    "tmp_used": 538333184,  
    "used_memory_samples": 6,  
    "cpu_user_time": 170,  
    "tx_bytes": 42,  
    "rx_bytes": 42,
```



Most recent 1000 invocations (8)
View performance or application logs of this function, or select a few invocations to view logs for selected invocations.

	Timestamp	Request ID	Trace	Error	Init duration	Duration	Memory %	CPU time	Network IO
<input type="checkbox"/>	2021-05-21 01:44:59 (UTC-04:00)	e1bda395-48bf-...	-	-	-	4ms	<div style="width: 37%;">37%</div>	50ms	1281 bytes
<input type="checkbox"/>	2021-05-21 01:39:23 (UTC-04:00)	c07d588b-c9ce-...	-	-	-	4ms	<div style="width: 37%;">37%</div>	10ms	6638 bytes
<input type="checkbox"/>	2021-05-21 01:39:23 (UTC-04:00)	38f00008-63e2-...	-	-	-	3ms	<div style="width: 37%;">37%</div>	10ms	-
<input type="checkbox"/>	2021-05-21 01:39:23 (UTC-04:00)	4d4386d6-b08b...	-	-	-	3ms	<div style="width: 37%;">37%</div>	20ms	576 bytes
<input type="checkbox"/>	2021-05-21 01:39:22 (UTC-04:00)	f5fb5273-8df3-...	-	-	-	20ms	<div style="width: 37%;">37%</div>	10ms	4740 bytes
<input type="checkbox"/>	2021-05-21 01:39:22 (UTC-04:00)	6637b76b-679f...	-	-	-	4ms	<div style="width: 37%;">37%</div>	10ms	3568 bytes
<input type="checkbox"/>	2021-05-21 01:39:21 (UTC-04:00)	7da05d90-41d5...	-	-	-	4ms	<div style="width: 37%;">37%</div>	-	3333 bytes
<input type="checkbox"/>	2021-05-21 01:39:21 (UTC-04:00)	17d0e85a-cd26...	-	-	-	16ms	<div style="width: 37%;">37%</div>	-	10558 bytes

1h 3h 12h 1d 3d 1w Custom  Add to dashboard  

Performance monitoring

Multi-function

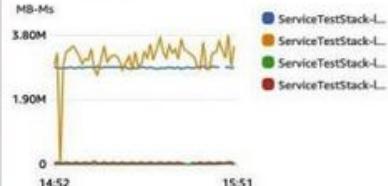
Filter metrics by function name

ServiceTestStack-lambda0LambdaFileSystemD0DC67D0-1K4PYS9EVIDAF  ServiceTestStack-lambdaMainLambdaMain6A167A80-1LWF0ES20UNL9 

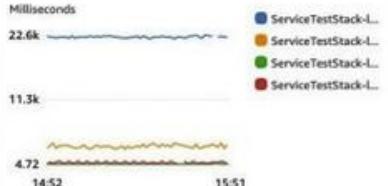
ServiceTestStack-lambdaMemoryLambdaMemory448FFFC-LUULBEPY0938  ServiceTestStack-lambdaCPULambdaCPU8B542329-8YUJON53M25E 

 In alarm 0  Insufficient data 0  OK 0

Function Cost



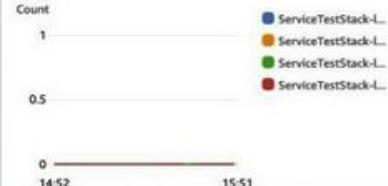
Duration



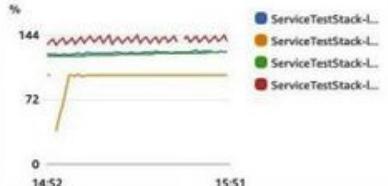
Invocations



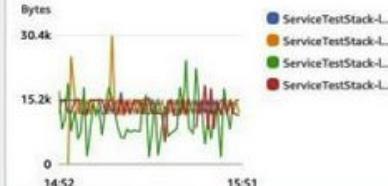
Errors



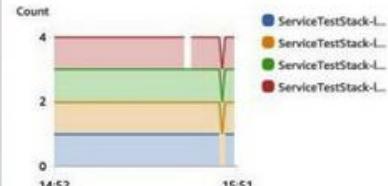
Memory Usage (Max)



Network Usage



Concurrent Executions Maximum



Throttles



Performance Optimization



How can we pick the right value?

Option 1: Trial and Error

Option 2: Lambda Power Tuning

 AWS Lambda Power Tuner

AWS Lambda Power Tuning is an open-source tool that can help you visualize and fine-tune the memory/power configuration of Lambda functions. It runs in your own AWS account - powered by AWS Step Functions - and it supports three optimization strategies: cost, speed, and balanced. Read more about it [here](#).

Lambda Power Tuner

What operation do you want to run?

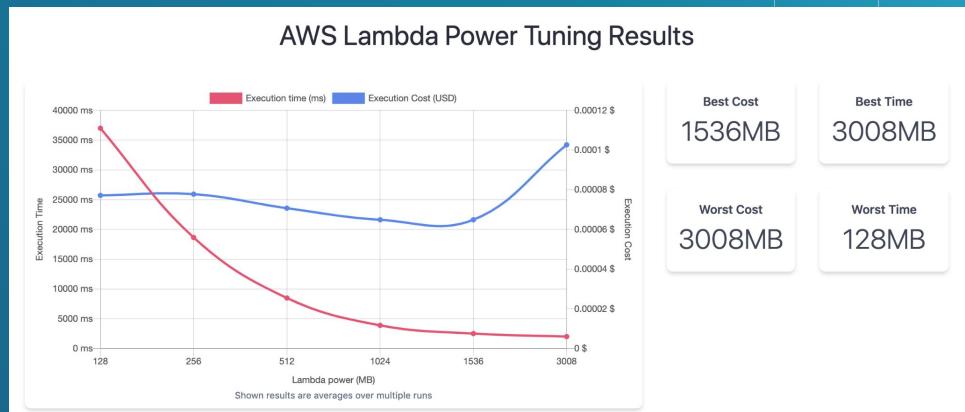
Lambda ARN
arn:aws:lambda:us-east-1:....

Strategy
 0.5 Invocation Number

Power values (MB)
 All possible values Custom values

Execute in Parallel
 Yes

Include JSON payload
 Yes





Logging - Hands On Example



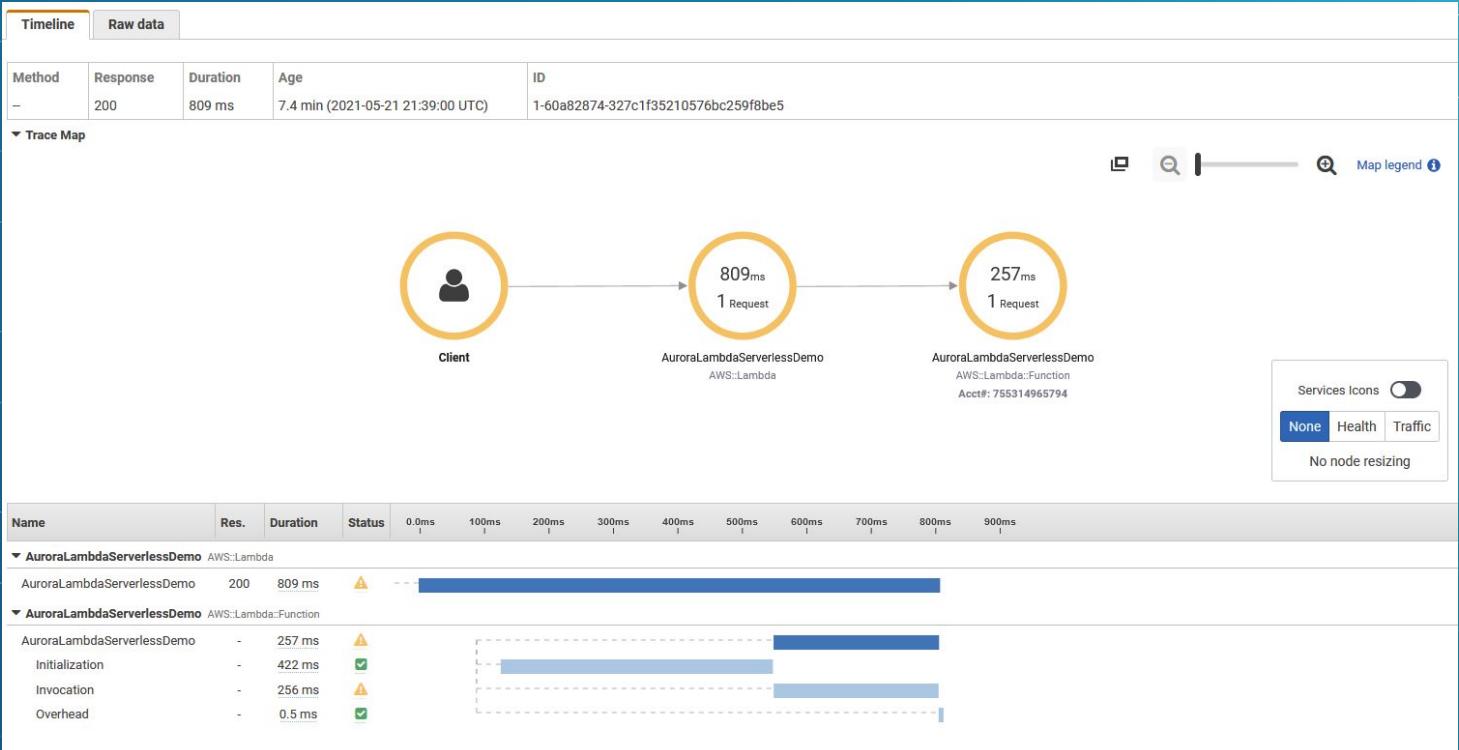
Monitoring - X Ray Tracing

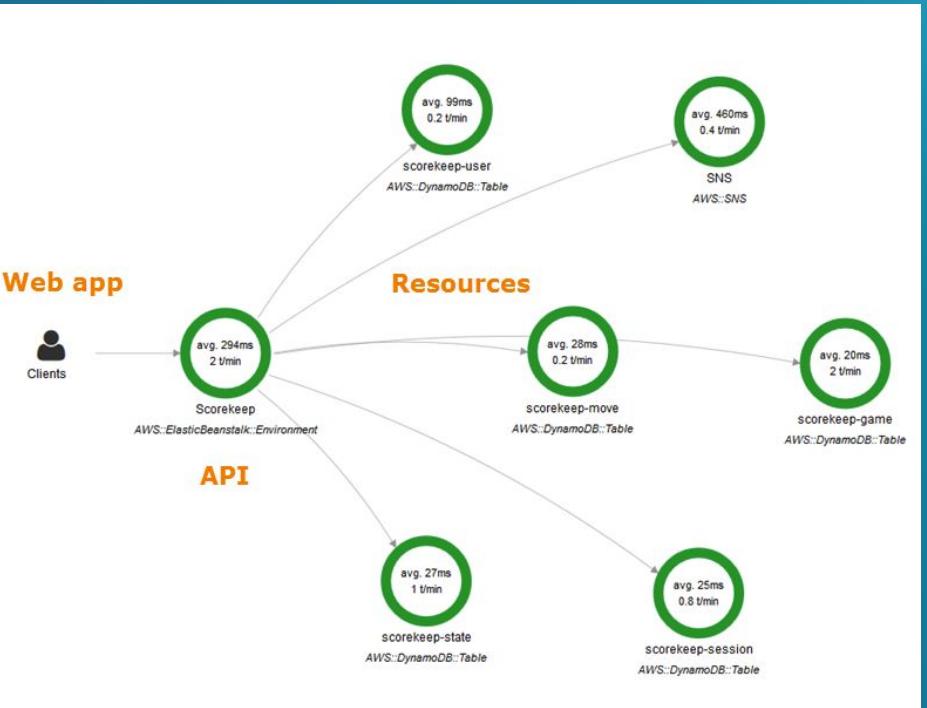


What is X Ray Tracing?

- ◊ Application level debugging feature
- ◊ Generates Trace information on your invocations
- ◊ Visualization aid for understanding how much time is spent where

You can add custom *segments* to your code for function traces







Pro Tips

- ◊ Its expensive
 - ◊ \$5.00 / 1 million traces **recorded**
 - ◊ \$0.50 / 1 million traces **retrieved**
 - ◊ \$0.50 / 1 million traces **scanned**
- ◊ More useful for ad-hoc debugging

Consider using sampling



Advanced Concepts - DLQ



What is a DLQ?

- ◊ Dead Letter Queue
 - ◊ Sometimes dependencies fail and messages cannot be processed...
 - ◊ DLQ is a destination to temporarily hold failed invocation message content
- After a number of attempts, messages can be sent to the DLQ for you to act on

Example

Method 1



Method 2





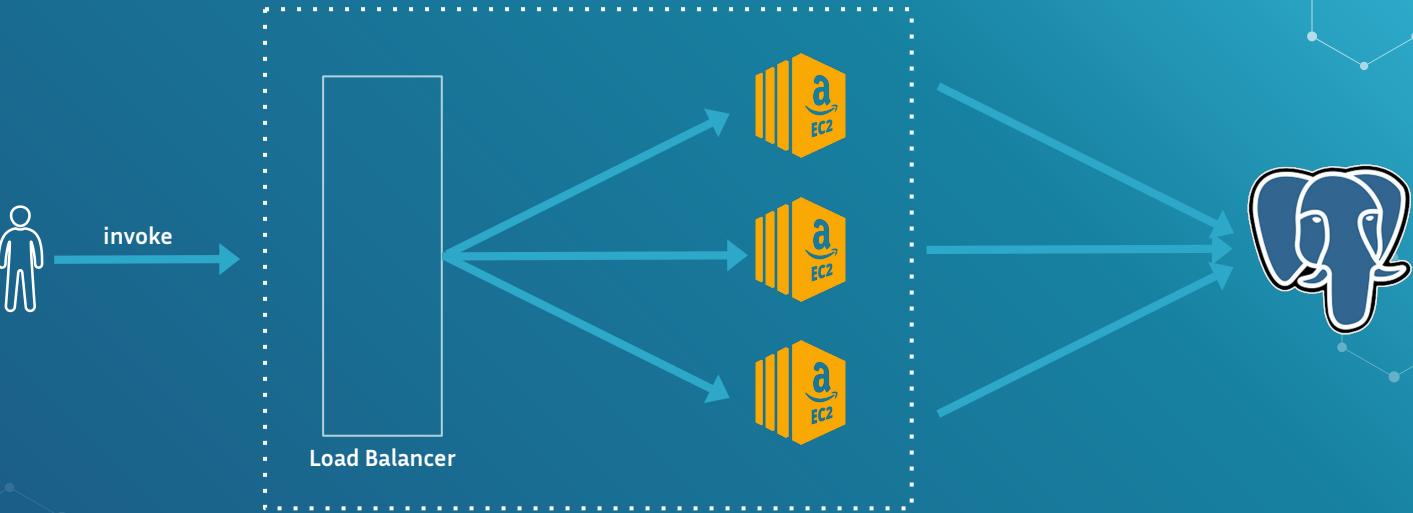
Best Practices

- ◊ Always use a DLQ for production applications
- ◊ Alarm on DLQ size > 0
- ◊ Be Proactive - have a process for re-driving messages



Advanced Concepts - RDS Proxy

Remember This?



Remember This?

invoke



RDS Proxy Feature

- ◊ Connection Pool Manager for RDS that integrates directly with Lambda
- ◊ Can be used with RDS MySQL, Postgres or Aurora
- ◊ If deciding to *not* use this feature, be careful with connection pool management



Good to Knows - Lambda Layers and Filesystem



Lambda Layers

- ◆ Allows you to easily add external libraries, data, config files to your function
- ◆ Zip the package contents and upload via console
- ◆ Only usable when uploading your function code as a .zip, and only 5 per function



Lambda Filesystem

- ◆ You can mount an Elastic File System volume to make its contents accessible to your function

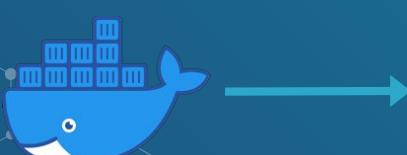


Good to Knows - Lambda & Docker



Lambda and Docker

- ◆ As of December 2020, you can deploy your functions using a Docker image
- ◆ 10GB size limit
- ◆ Base images are provided





Good to Knows - Continuous Deployment



Lambda CI / CD

- ❖ For larger projects, you may want to consider CI / CD pipelines
- ❖ Integrate with Github/CodeCommit and Codebuild for an automatic build system
- ❖ Leverage CodePipeline to orchestrate the changes and update your function

Example CI/CD Pipeline



buildspec.yml



AWS SAM



AWS CDK



AWS CodeBuild



AWS CodePipeline



AWS CloudFormation



AWS Lambda





Practical Content - Overview



What's Next?

- ❖ Pre-requisites: AWS CLI, CDK, PIP
- ❖ **Console:**
 - ❖ Cloudwatch Event + Lambda
 - ❖ HTTP API + Lambda

CDK:

- ❖ SQS + Lambda